

# AI-Driven Analysis of BC Post-Secondary Institution Strategy Documents

---

Helping consultants and policymakers extract strategic insights  
from institutional documents using AI

# Current Scope (MVP Focus)



Analyzing institutional strategies manually is time-consuming. Our AI solution automates insight extraction



To build a focused and high-impact MVP, we targeted four key PDF documents per institution:

Strategic Plan

Financial Statement

Government Mandate Letter

Courses List



Total: 40 documents across 10 institutions




Each document was parsed page-by-page and tagged with standardized metadata for institution name, document type, and source



**Decision-makers** can identify trends across institutions through an interactive user interface

# Progress



PHASE 1 – INGESTION  
PIPELINE




PHASE 2 – BUILD  
KNOWLEDGE GRAPH



PHASE 3 – SEMANTIC  
SEARCH &  
EMBEDDING



PHASE 4 –  
CONSULTANT-FACING  
SEMANTIC INTERFACE



# Phase 1 – Ingestion Pipeline (Completed)

PDFs loaded using [LangChain's PyPDFLoader](#)

Each page treated as an individual content unit

Custom metadata tagging applied:

- institution
- source\_file
- doc\_type
- content\_type

documents.json – full text and metadata per page

pdf\_metadata.json – summary metadata only

*Low-value pages are skipped using a hybrid approach: text density thresholds and keyword-based filtering, balancing **efficiency and accuracy** while ensuring valuable content isn't lost.*

# Sample Data Snapshot

```
{  
  "content": "...Strategic goals include...",  
  "metadata": {  
    "institution": "Langara College",  
    "source_file": "Langara Strategic Plan.pdf",  
    "doc_type": "Strategic Plan",  
    "content_type": "PDF"  
  }  
}
```

- Clean and structured for downstream use
- Enables fast search, QA, or filtering
- Supports graph and semantic search applications

# Phase 2 – Knowledge Graph (In Progress)

## Goal:

- Extract (subject, predicate, object) triples from raw document text using OpenAI GPT3.5 Turbo

## Examples:


- Selkirk College, offers, Renewable Energy Diploma
- BCIT, partnered\_with, Indigenous Communities

## Tools:

- spaCy for named entity extraction (NER)
- OpenAI / OpenRouter LLMs for triple generation
- Output saved as triples.jsonl

## Quality Assurance:

- evaluating knowledge graph accuracy by comparing extracted triples against a reference dataset



# LLM Pipeline (In Progress)



We use OpenAI GPT3.5 Turbo for relationship extraction



Each page is chunked (up to ~2000 characters) and sent to the LLM with structured prompts



Extracted triples are returned in JSON format and tagged with source metadata



Deloitte has provided \$62 worth of OpenRouter tokens, but access was unsuccessful for unknown security reasons, so OpenAI token used instead.



Efficient use of prompts is required to stay within budget

Skipping short or low-value pages

# Phase 3 – Semantic Search & Embedding (Next)

## Planned Next Steps:

- Chunk pages into ~500-token segments
- Embed using OpenAI or Mistral models
- Store in FAISS or Pinecone vector database

## FAISS vs. Pinecone:

- FAISS for local efficiency
- Pinecone for managed scalability

## Purpose:

- Enable consultants to ask natural questions like:
- *“Which colleges offer AI-related diplomas?”*
- *“What sustainability goals exist across institutions?”*



# Phase 4 – Consultant-Facing Semantic Interface



## **Streamlit-Powered QA Interface:**

Simple UI to type questions about BC institutions

LLM-backed retrieval pipeline fetches best answers

Includes citations to source PDFs and institutions

Includes a simple feedback mechanism (e.g., thumbs-up/down for answers)



## **Benefits:**

Fast, explainable answers

Document traceability

Strategic insight engine for RFPs and client work

# Technical Stack Summary

- Tools Used So Far:
  - LangChain (PDF loading, RAG pipeline)
  - spaCy (Named Entity Recognition)
  - OpenRouter (LLM access)
  - Python 3.10 (venv, CLI)
  - JSON (document storage, triplet output)
  - Ready for Neo4j, FAISS, or Streamlit to be decided for UI

