Deloitte's Three-Agent AI Architecture: Strategic Insights for BC Post-Secondary Institutions

Assisting Deloitte consultants in extracting strategic insights from their latest institutional documents while minimizing hallucination.

The Problems

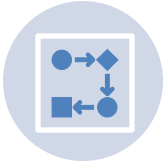

The research process is time-consuming

The existing AI tool is too generalized to provide tailored responses

Hallucination exists in AI-generated answers

GPT not scalable for large document sets

No persistent retrieval or memory beyond the session

Why is our framework better than just using an OpenAI API?

We choose the embedding model, optimizing for accuracy and cost.

We define how documents are chunked, structured, and refined — which dramatically improves answer quality.

We can scale to thousands of documents and handle complex logic across multiple data types.

We can integrate semantic triples, institutional metadata, and advanced filtering — something no OpenAI assistant can currently do.

Project Workflow Overview

1

PHASE 1 – DATA INGESTION PIPELINE

2

PHASE 2 – BUILD KNOWLEDGE GRAPH

3

PHASE 3 – SEMANTIC SEARCH & EMBEDDING

4

PHASE 4 – CONSULTANT-FACING SEMANTIC INTERFACE

Pinecone-based Question Answering System Overview

System Components

-Pinecone vector database for similarity search

-OpenAI API for embeddings and answer generation

-Query processor handling user interactions

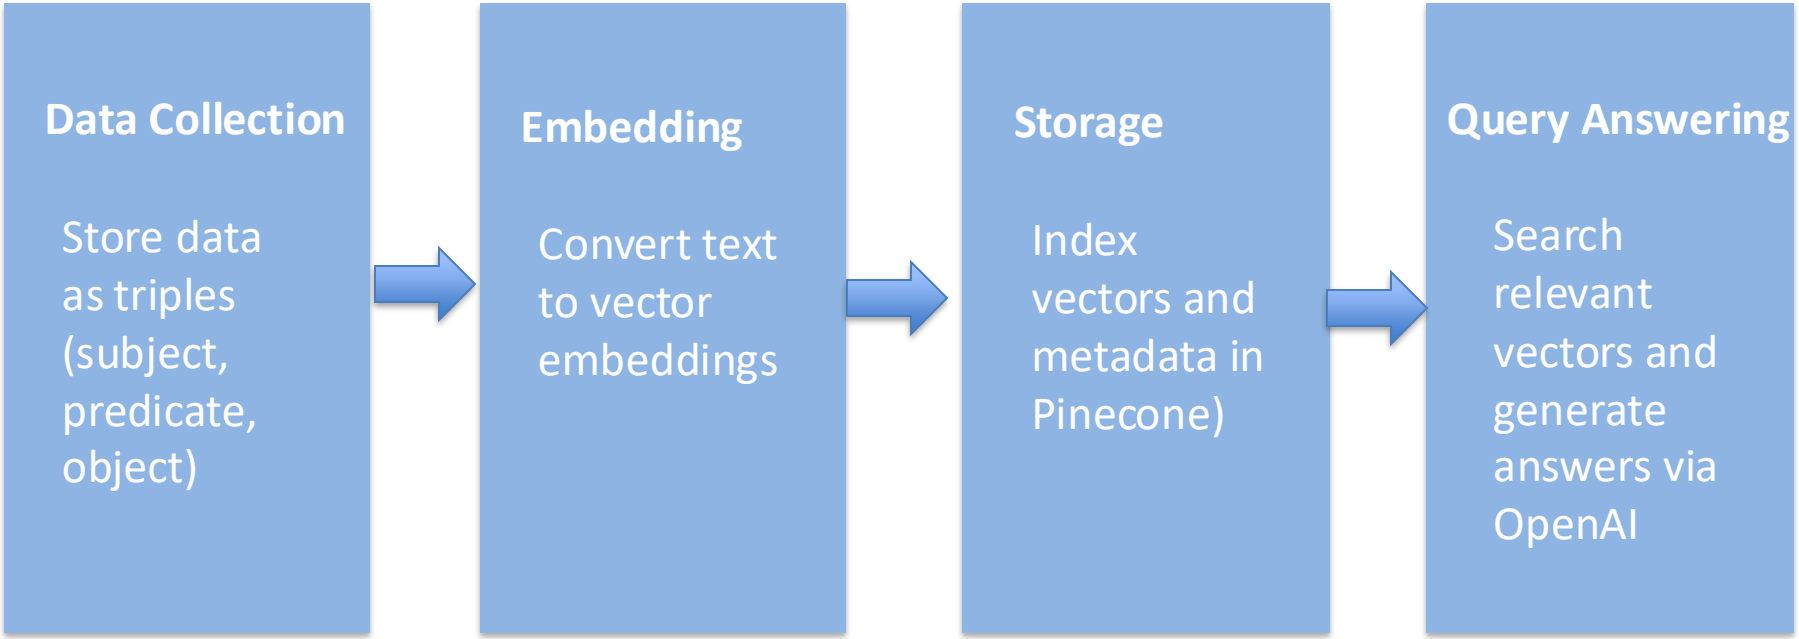

Using Pinecone Vector Database

Database Setup
Connect via Pinecone API key. Create and manage indexes for vector storage.

Text Embedding
Convert text into vectors using OpenAI models. The embed_text() function transforms text into numerical representations.

Query Search
Convert questions into vectors. Search for similar vectors in Pinecone to retrieve relevant information.

Query Answering System Workflow

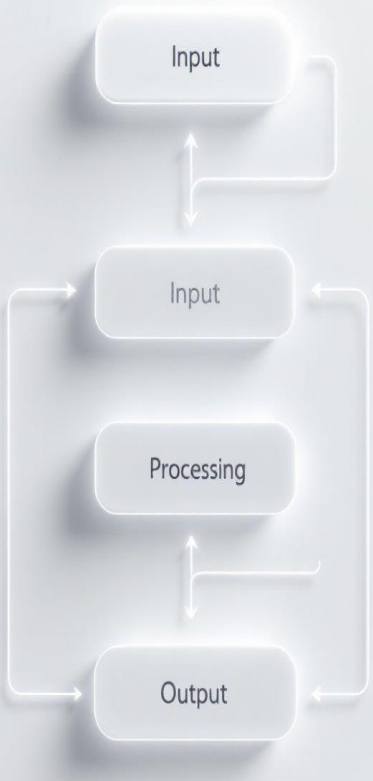

User Query Input
User enters a question like "Tell me about research achievements."

Query Embedding Generation
The get_query_embedding() function converts text to vector format.

Pinecone Similarity Search
The search_pinecone() function finds relevant information triples.

OpenAI Answer Generation
The ask_openai() function sends relevant context to GPT-3 for answer generation.

AI-Powered User Interface

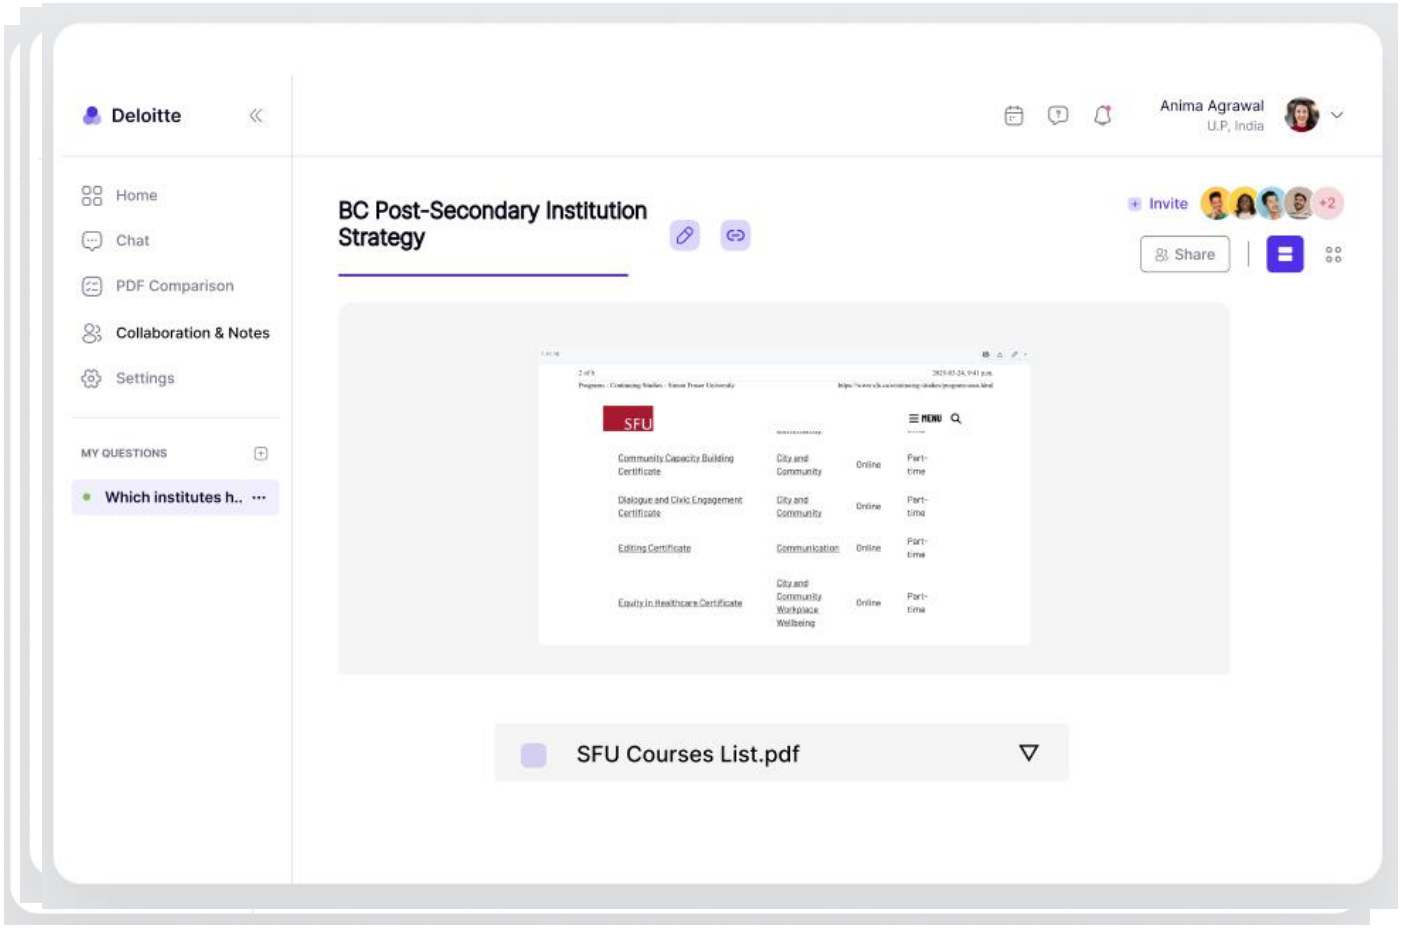

Potential Use Cases for Deloitte Consultants

Rapid benchmarking of institutional goals, programs, and metrics

Identification of curriculum sharing or partnership opportunities

Automated extraction of KPI-relevant insights from mandate letters and strategic plans

Accelerated response to RFPs by surfacing aligned institutional capabilities

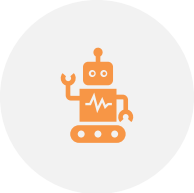

Detection of trends across institutions (e.g., sustainability, AI, Indigenous initiatives)

Future Advancements

- Connet to KX to leverage its context and make AI suggestions
- Connect to Omnia to leverage its financial analysis features
- Adaptive AI Learning & Customization
- Multi-Language & Cross-Region Support
- Expanded Data Sources & Integrations

Thank you